

5LIA0 - Embedded Visual Control

Computer Vision



Egor
Bondarev

- MSc in Robotics, Belarus
- Invention Machine Corp – Artificial Intellect
- OOTI, Software Technology, Post-master, TUE
- PhD on Performance of Real-Time Systems, TUE
- Currently:
 - Assistant Professor at TUE: 3D Sensing and Reconstruction
 - CEO of IVRA Holding

VCA Group Experience: Sensing Platforms

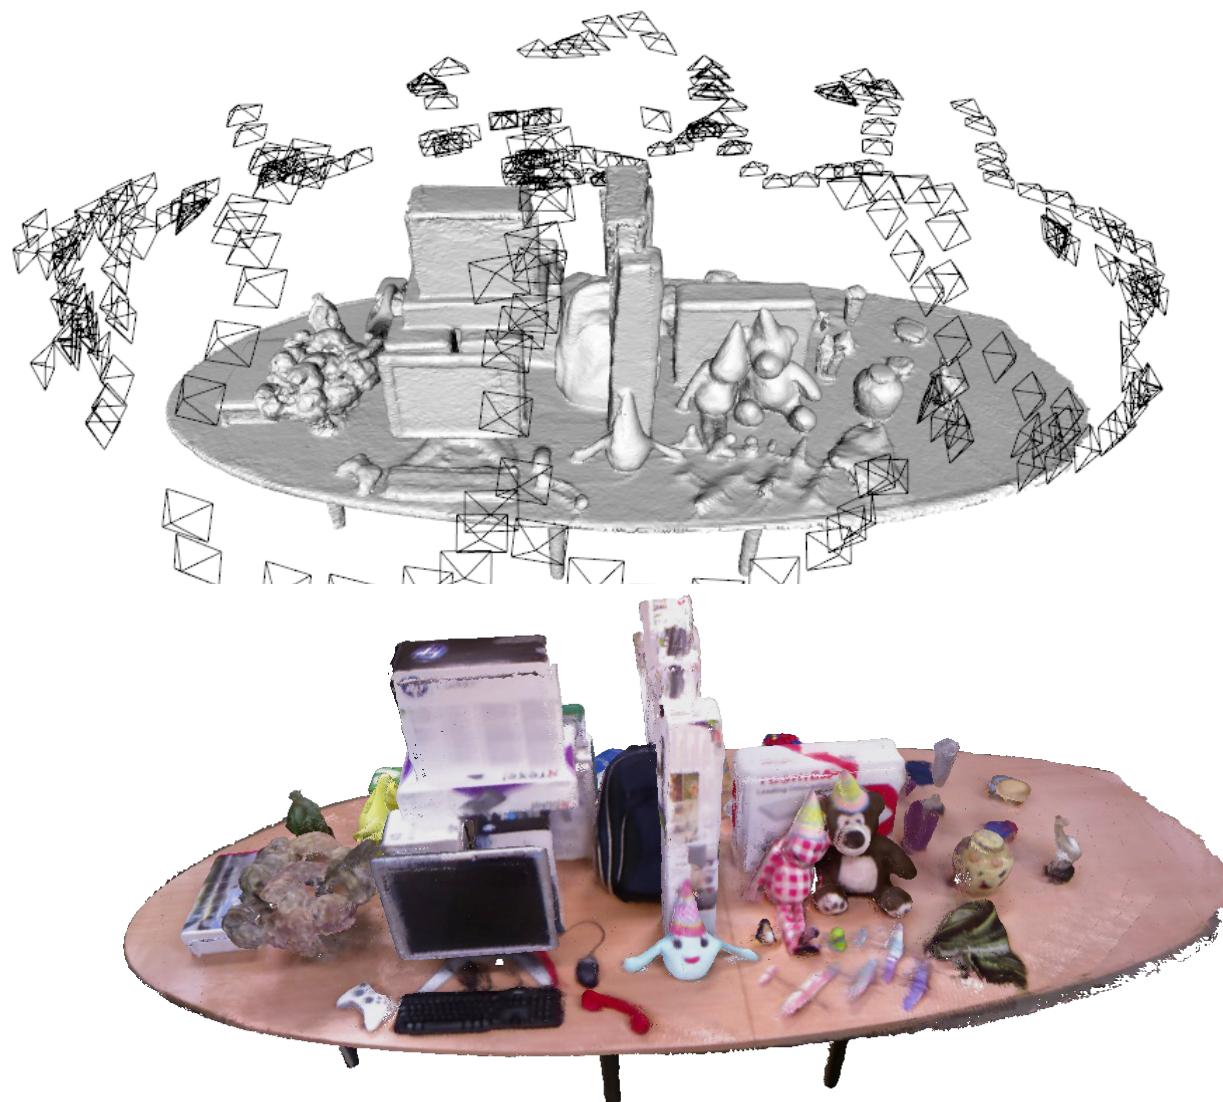
Set of 2D/3D Sensing Platforms

- Few laser-scanner platforms (LIDARs)
- RGB-D mapping system (Kinect)
- Stereo Cameras, Panoramic 360 camera
- Autonomous robot
- Thermal cameras

All of them have 1 or more computers inside



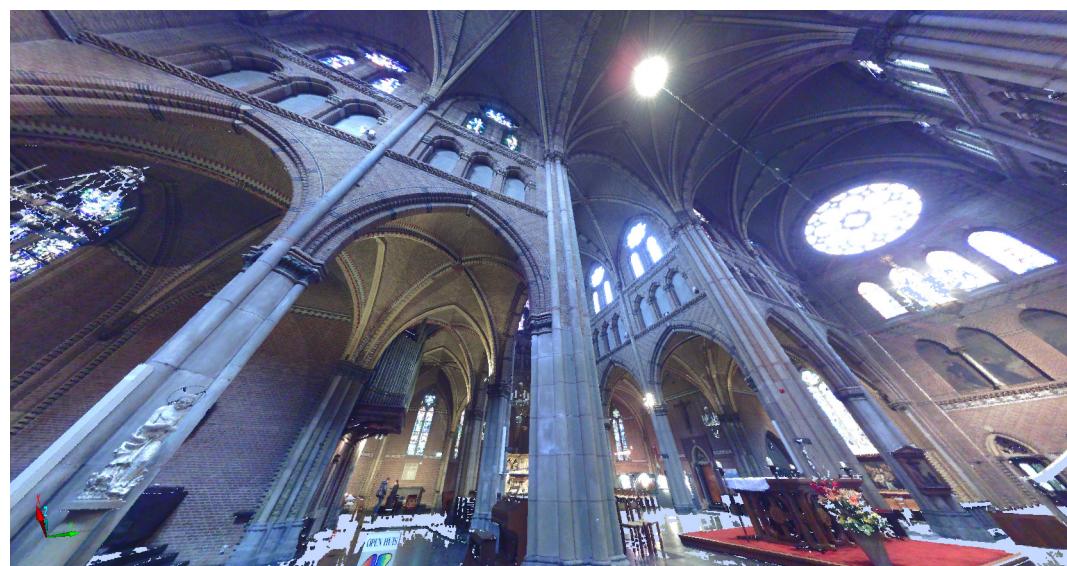
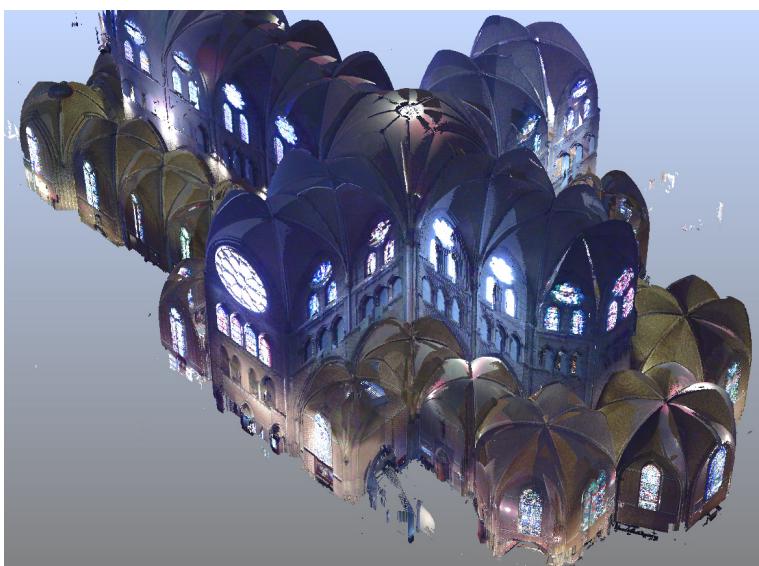
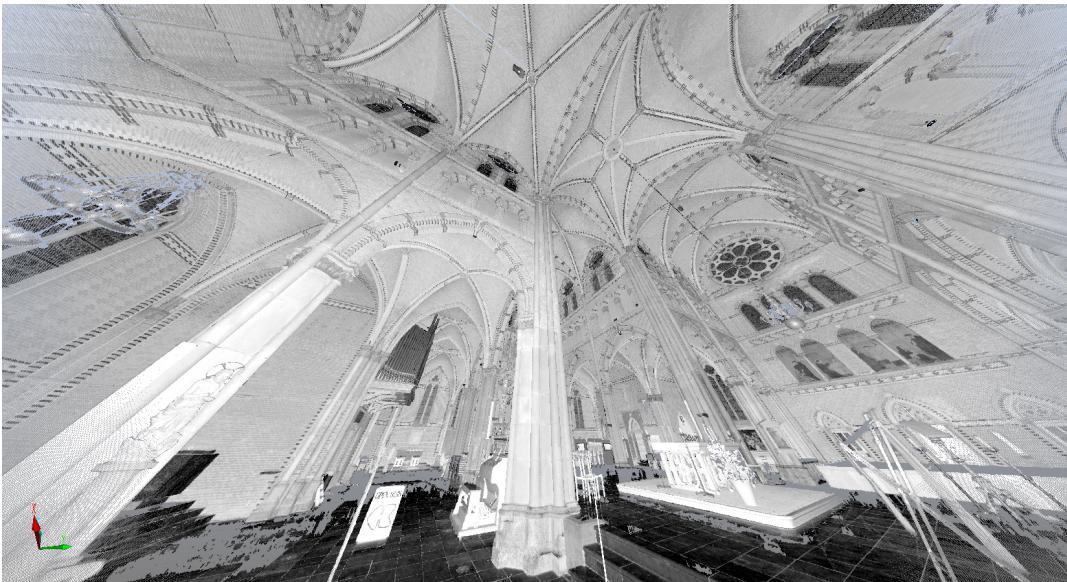
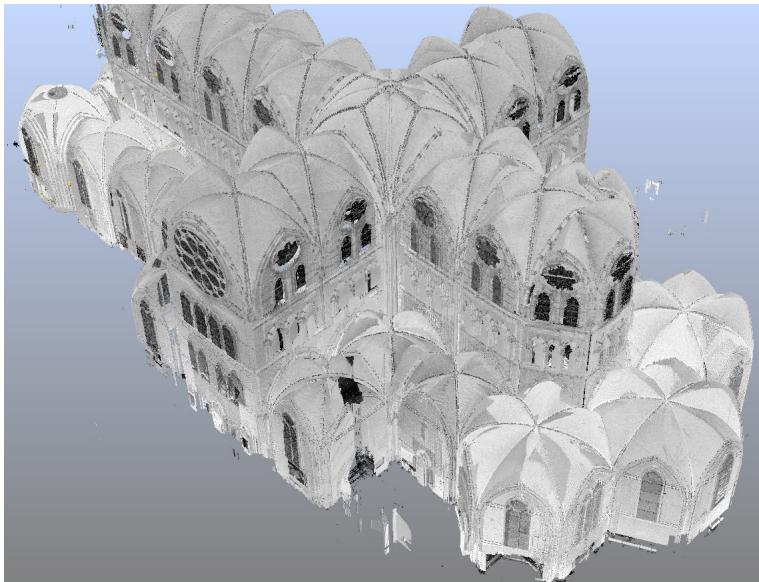
What we do at VCA: Detailed 3D Reconstruction & Texturing



What we do at VCA: Complex spaces into 3D



What we do at VCA: Extremely Large Spaces

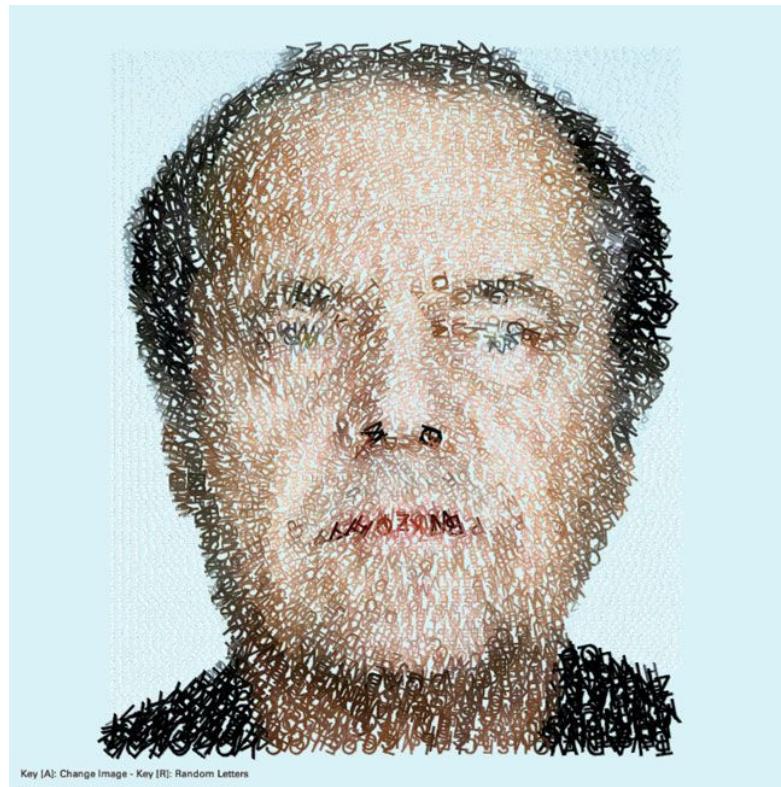


What we do at VCA: Demo. Room in 3D

These hours:

- Computer Vision Introduction
- Vision for Autonomous Vehicles
- Lane Detection
 - Canny filter
 - Hough Transform
- Traffic Signs Detection
 - Next lecture

Computer Vision in Genera



Key [A]: Change Image - Key [R]: Random Letters

Make computers understand images and videos



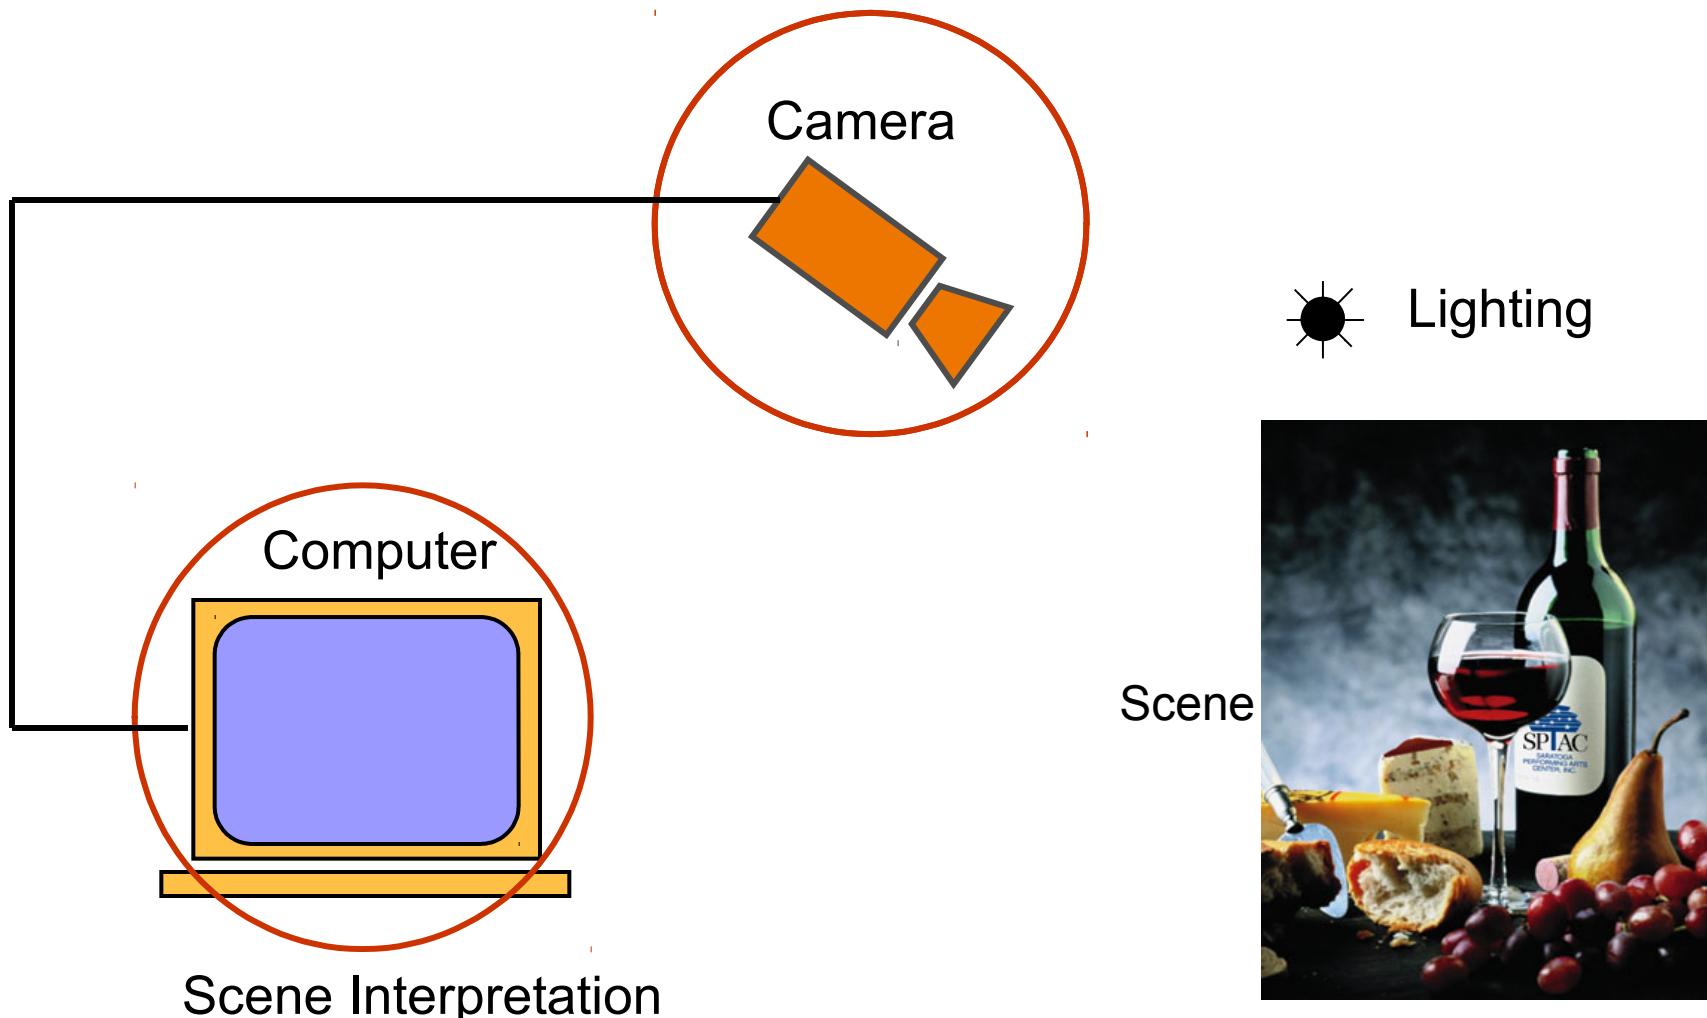
What kind of scene?

Where the camera is located at this scene?

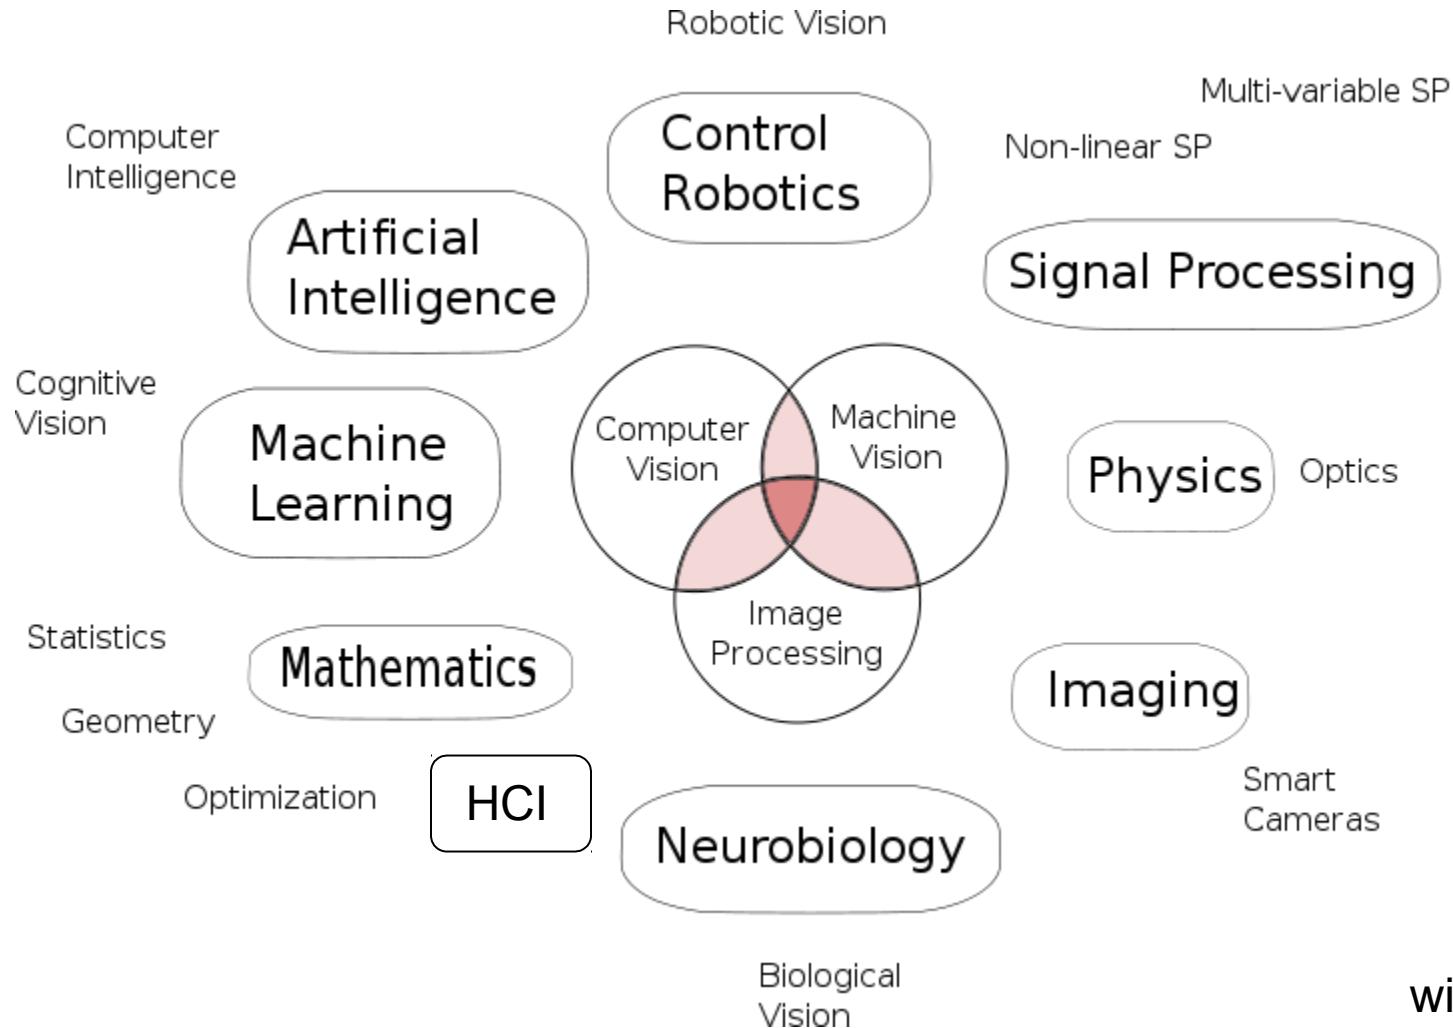
Where are the cars?

How far is the building?

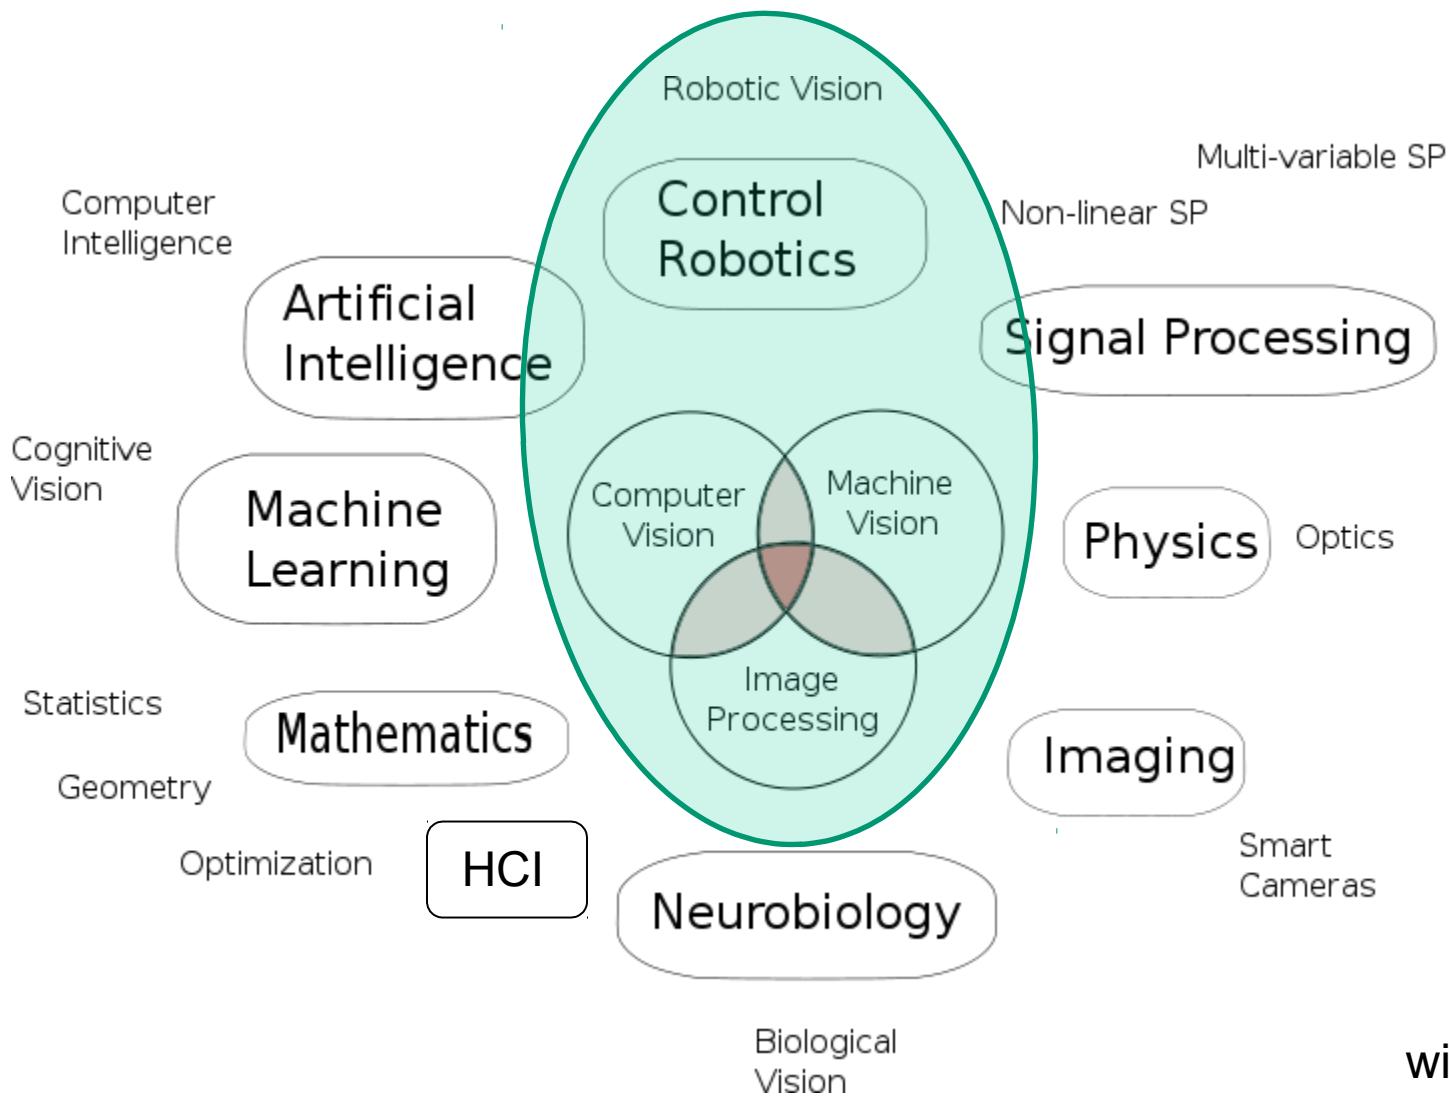
Components of a computer vision system



Application Domains



Application for Robotics



Computer Vision for Autonomous Vehicles

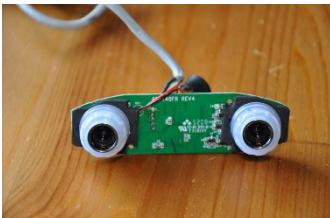


- Major goals
 - Identify camera (car) *position and orientation* at every moment in time - localization
 - Understand surroundings – mapping of scene around
 - Detect and recognize objects:
 - Traffic signs
 - Pedestrians
 - Lanes
 - Other vehicles
 - Track (to analyze situation around _):
 - Road bumps – to control suspension
 - Cars around – for collision avoidance, platooning
 - Compute optimal path from A to B
- Ultimate goal
 - Enable fully autonomous driving

RGB camera



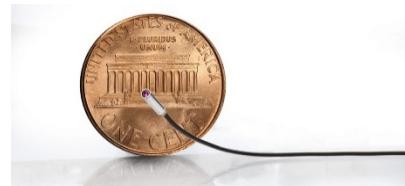
Stereo camera



IR sensors



Ultrasound sensors



Thermal camera



Laser scanners - LiDAR



Supportive sensors:

GPS,

Odometers,

IMU

(accelerators,

gyroscopes,

magnetometers,

altimeters),

Acoustic

- RGB camera
 - Advantages
 - Low cost: from 1 euro
 - Long range: till horizon
 - Lots of rgb-based image processing algorithms available
 - High resolution
 - Data post-processing is not required
 - Disadvantages
 - No depth data, 2D picture
 - Low performance for occlusions, cluttered scenes
 - Blind at night, low performance in fog, rain
 - High data rate
 - For which functions
 - All of the described above



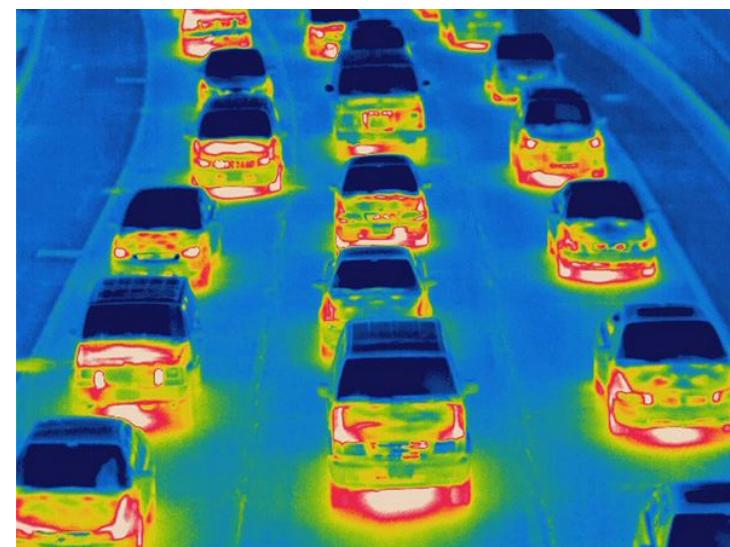
- Stereo camera
 - Advantages
 - Depth data can be obtained!
 - High performance for occlusions, cluttered scenes
 - Long range: till horizon, depth range is limited
 - Depth-based image processing algorithms available, but limited
 - Relatively high resolution (HD)
 - Disadvantages
 - High cost: from 100 euro
 - Data post-processing is required: disparity map computation is heavy
 - Blind at night, low performance in fog, rain
 - Noisy depth data
 - For which functions
 - All *short and medium* range functions described above



- Kinect-like (Structured Light, or Time of Flight)
 - Advantages
 - Depth data is directly available!
 - Data post-processing is not required
 - Can see depth at night
 - High performance for occlusions, cluttered scenes
 - Depth-based image processing algorithms available, but limited
 - Relatively high resolution (HD)
 - Disadvantages
 - Very short range: up to 50 m
 - High cost: from 100 euro
 - Does not work outdoor!
 - Noisy depth data
 - For which functions
 - Only *short* range functions described above



- Thermal camera
 - Advantages
 - Can see perfectly at night
 - Long range
 - Data post-processing is not required
 - Low data rate
 - Disadvantages
 - High cost: from 300 euro
 - Low resolution
 - For which functions
 - Mostly for night vision, human and car detection in cluttered scenes



- IR, Ultrasound sensors
 - Advantages
 - Low cost
 - Robust, works at night, rain, fog
 - Disadvantages
 - Short range: up to 20 m
 - Very low resolution
 - For which functions
 - Used mostly for obstacle detection

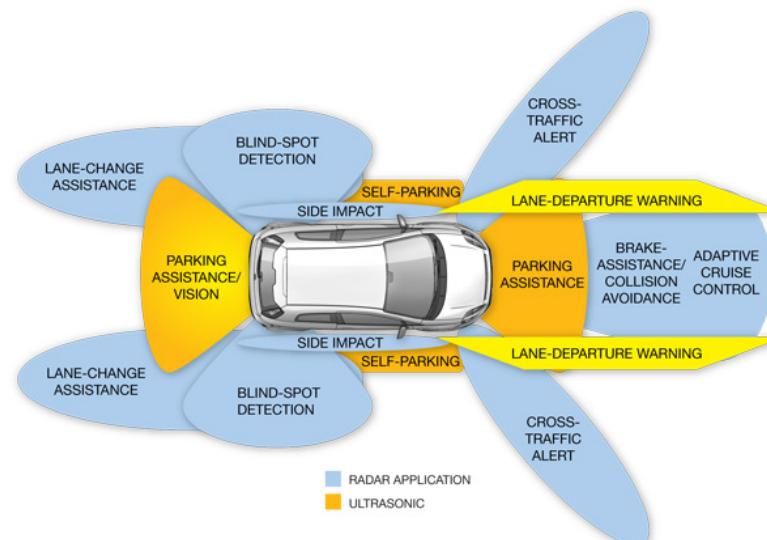
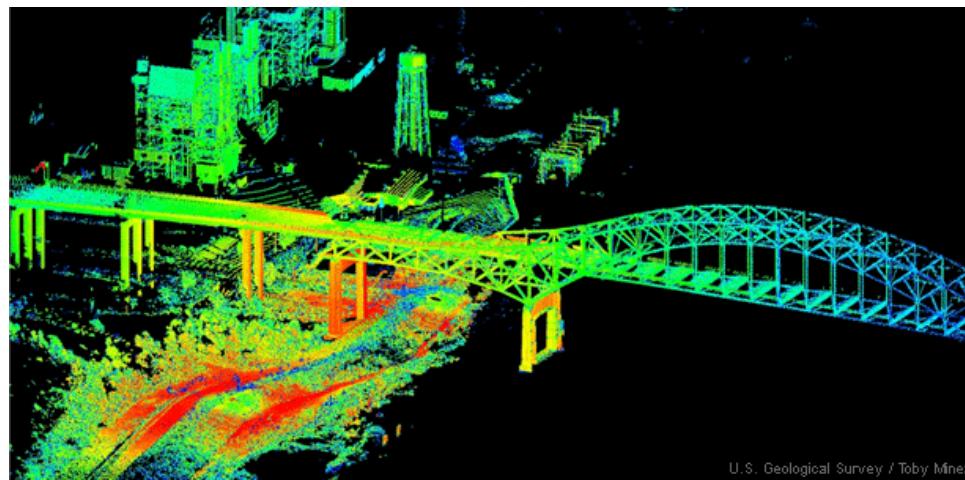


Figure 2 Several driver-assistance systems are currently using radar technology to provide blind-spot detection, parking assistance, collision avoidance, and other driver aids (courtesy Analog Devices).

- Laser scanner - LiDAR
 - Advantages
 - Full 3D data around
 - Point-clouds algorithms available, limited
 - High performance for occlusions, cluttered scenes
 - Very high resolution, 10 mln points/second
 - Data post-processing is not required
 - Can see at night, fog and rain
 - Disadvantages
 - Very high cost: 6000 euro
 - Medium range: up to 300 meters
 - Very high data rate
 - For which functions
 - All of the described above + 3D reconstruction of road map

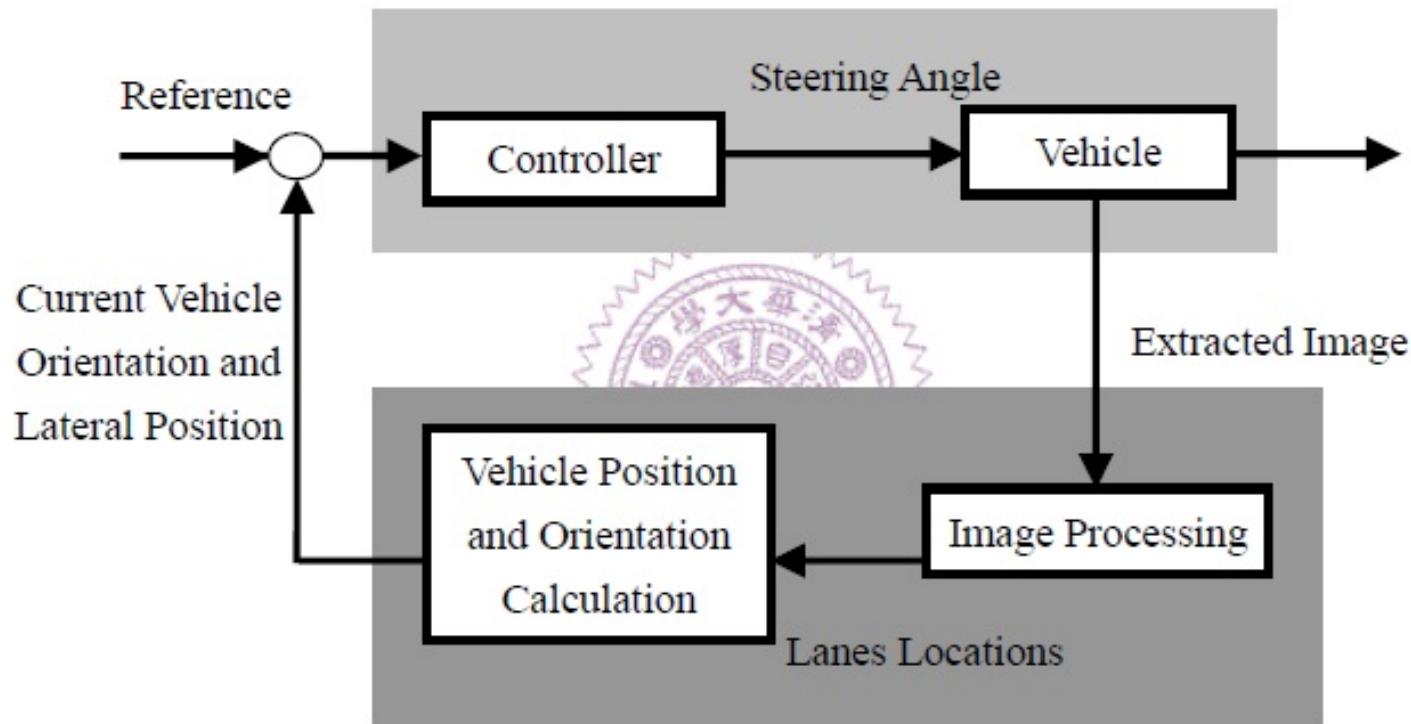


Specific Vision Tasks

Lane Tracking



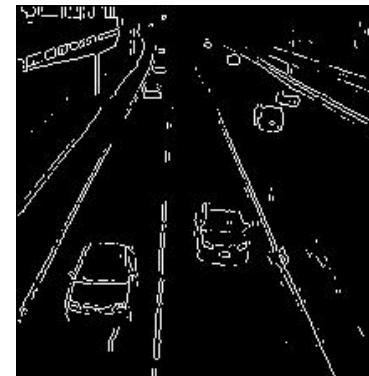
- Challenging task
 - Varying road conditions vary
 - Occlusions (parked vehicles)
 - Lane curves
 - Shadows
 - Slopes
 - Tough weather conditions
- Two types of approaches
 - Model based
 - Less sensitive to lane quality, occlusions
 - Not robust for unexpected scenes and non-standard lane marking
 - Computationally expensive – model fitting optimization
 - Feature based
 - Very sensitive to occlusions, poor road marking
 - Robust to an arbitrary marking
 - Less computationally expensive



Conventional Lane Detection Pipeline (many others exist)

- Step 0: Pre-filtering: subtract background, filter image
- Step 1: Get the edge information
- Step 2: Hough Transform
- Step 3: Search out the lane marking candidates
- Step 4: Decide the lane marking

- Step 1: Get the edge information.
 - Use for example Canny detector. OpenCV source code:
[http://
docs.opencv.org/2.4/doc/tutorials/imgproc/imgtrans/canny
_detector/canny_detector.html](http://docs.opencv.org/2.4/doc/tutorials/imgproc/imgtrans/canny_detector/canny_detector.html)

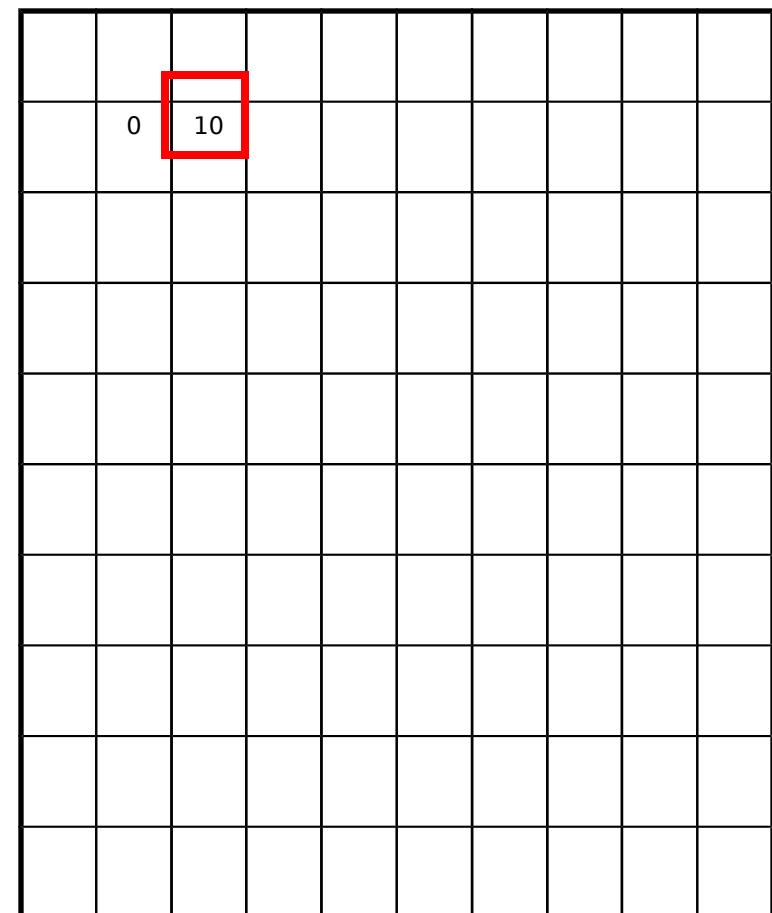


1. Filter image to remove noise – Gaussian filter
2. Find the intensity gradient of the image
3. Non-maxima pixels suppression
4. Remove or leave each pixel based on hysteresis

1. Filter image to remove noise – Box Filter

$$g[\cdot, \cdot] \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0



1. Filter image to remove noise – Box Filter

$$g[\cdot, \cdot] \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

1. Filter image to remove noise – Gaussian Filter

Gaussian coefficients

0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

Mask 5 x 5

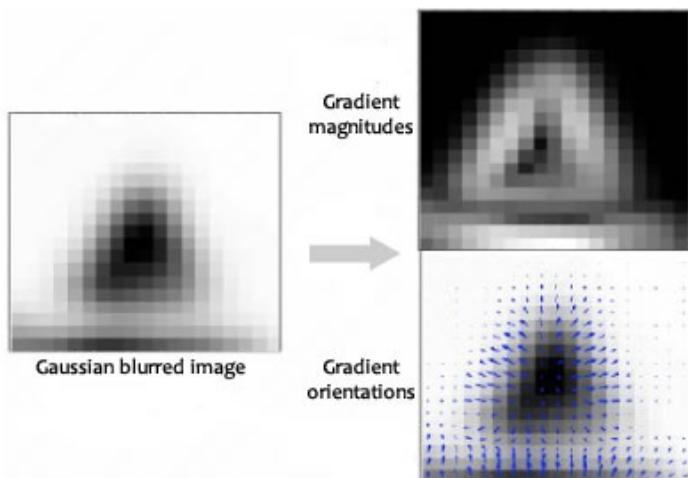
Instead of box filter

coefficients

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

32 Canny Detector

2. Find the intensity gradient of the image



2.1. Apply a pair of convolution masks in X and Y directions:

$$G_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix}$$

$$G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ +1 & +2 & +1 \end{bmatrix}$$

2.2. Find the gradient strength and direction with:

3. Non-maxima pixels suppression:

reducing thickness of the edges

by leaving only pixels with highest gradient

$$G = \sqrt{G_x^2 + G_y^2}$$
$$\theta = \arctan\left(\frac{G_y}{G_x}\right)$$



3. Remove or leave each pixel based on *hysteresis*

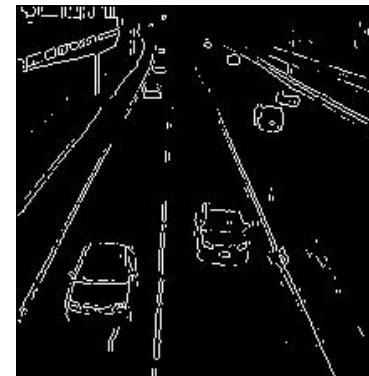
The final step. Canny uses two thresholds (upper and lower):

- If a pixel gradient is higher than the *upper* threshold, the pixel is accepted as an edge
- If a pixel gradient value is below the *lower* threshold, then it is rejected.
- If the pixel gradient is between the two thresholds, then it will be accepted only if it is connected to a pixel that is above the *upper* threshold.



Conventional Lane Detection Pipeline

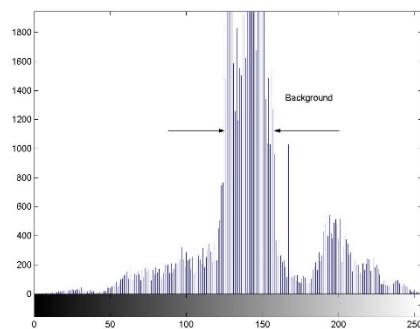
- Step 1: Get the edge information.
 - Use for example Canny detector. OpenCV source code:
[http://
docs.opencv.org/2.4/doc/tutorials/imgproc/imgtrans/canny
_detector/canny_detector.html](http://docs.opencv.org/2.4/doc/tutorials/imgproc/imgtrans/canny_detector/canny_detector.html)



Lost many edges



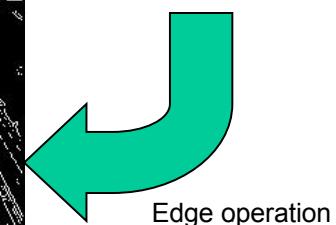
- Step 0: Filter the image.
 - Option: background subtraction



Use global histogram to find the background gray range and subtract it from the original image



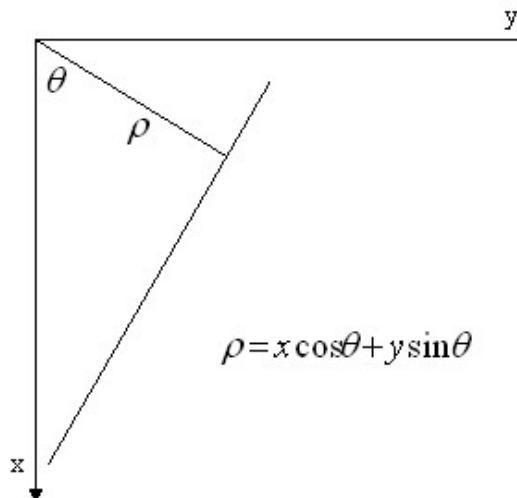
Compare these images
Edges are preserved using background subtraction method



Edge operation

- Step 2: Hough Transform

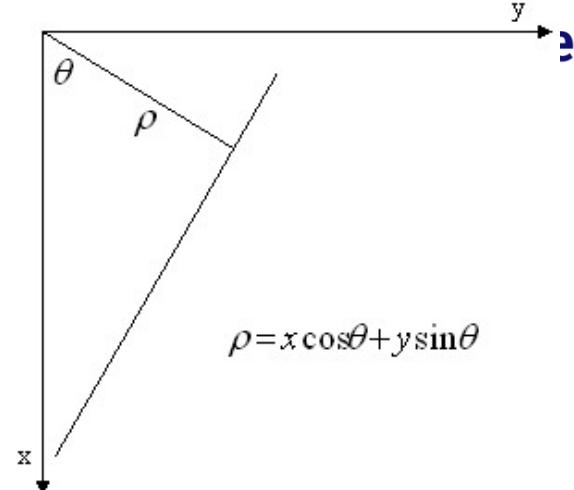
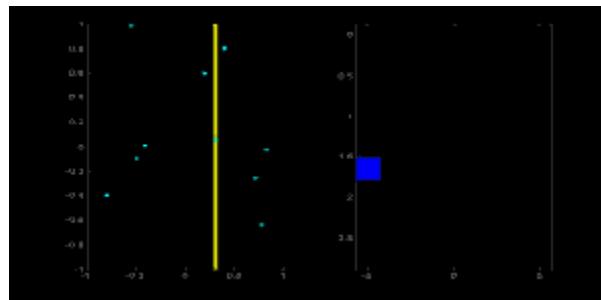
- technique to detect any shape, if you can represent that shape in mathematical form
 - line can be represented as $y = mx + c$ or



- where *rho* is the perpendicular distance from origin to the line, and *theta* is the angle

Conventional Lane Detection Pipeline

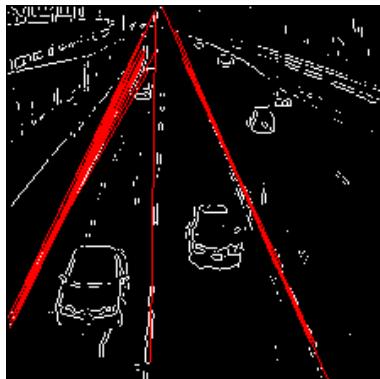
- Step 2: Hough Transform



- Hough Transform creates a 2D array [rho, theta] or accumulator (to hold values of two parameters) and it is set to 0 initially.
- Take the first Canny pixel. You know its (x,y) values. Now in the line equation, put the values \theta = 0, 1, 2, ..., 180 and check the \rho you get. Store every (\rho, \theta) pair into 2D array.
- Do the same for other pixels found by Canny detector.
- Count equal elements (clustering) in 2D array – these counts show how many pixels belong to each hypothetical line
- There will be a lot of non-clustered elements with 1 counts – no line. But some clusters will have from 2 to many counts – candidates for lines.

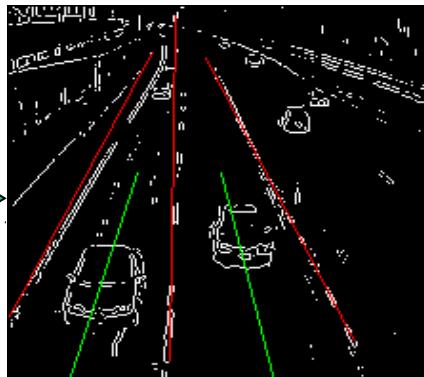
OpenCV source code: http://docs.opencv.org/3.0-beta/doc/py_tutorials/py_imgproc/py_houghlines/py_houghlines.html

- Step 3: Search out the lane marking candidates.

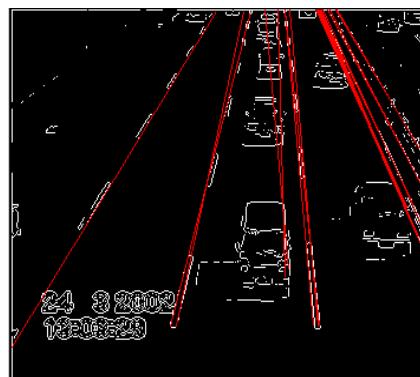


1. The red lines are the first 20 lines which have the biggest count numbers in Hough parameter space.
2. For each lane marking in the real scene, there are many line candidates around it.
3. There are some fake lines caused by the vehicle queue.

- Step 4: Decide on the lane marking



1. Sort the candidate lines by their position from left to right
2. Around each line cluster, choose the candidate which has the biggest count number as the lane marking in real scene
3. Delete the fake lane marking candidates
4. Calculate the mid-line of each lane (shown as green lines)

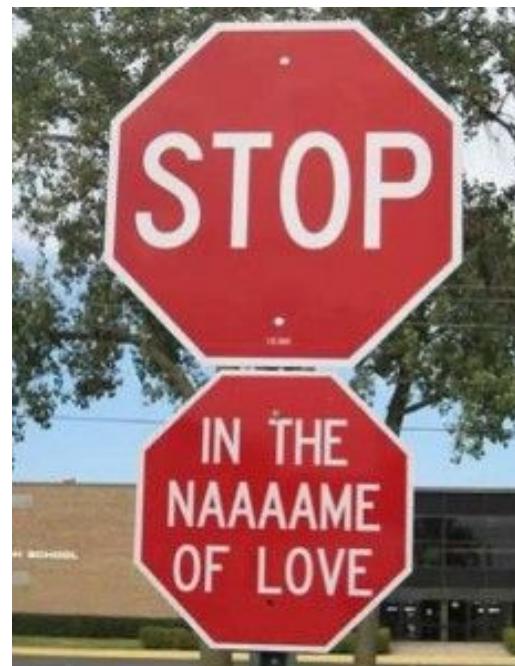


- Steps 5-100

- You may want to project those lines to top view of the road scene.
- Locate your vehicle on the this top view scene.
- Make any smart decisions, based on found lanes and your vehicle location.
- You may target more advanced cases (curved lanes, occlusions, bad road marking) by methods, like:
 - Deformable templates
 - Multi-resolution Hough Transform
 - B-snake
- Some references:
 - Karl Kluge, Sridhar Lakshmanan, “A deformable-template approach to lane detection”,
 - Gonzalez, J.P.; Ozguner, U.; “Lane detection using histogram-based segmentation and decision trees”,
 - Yu, B.; Jain, A.K.; “Lane boundary detection using a multiresolution Hough transform”,
 - Yue Wang; Eam Khwang Teoh; Dinggang Shen; “Lane detection using B-snake”
- Apart from OpenCV, some other open source libraries for lane detection:
 - <https://github.com/jdorweiler/lane-detection>
 - <http://www.vision.caltech.edu/malaa/software/research/caltech-lane-detection/>
 - <https://github.com/tomazas/opencv-lane-vehicle-track>

Specific Vision Tasks

Traffic Signs Detection



- Computer Vision Introduction
- Vision for Autonomous Vehicles
- Lane Detection
 - Canny filter
 - Hough Transform
- Traffic Signs Detection
 - Next lecture